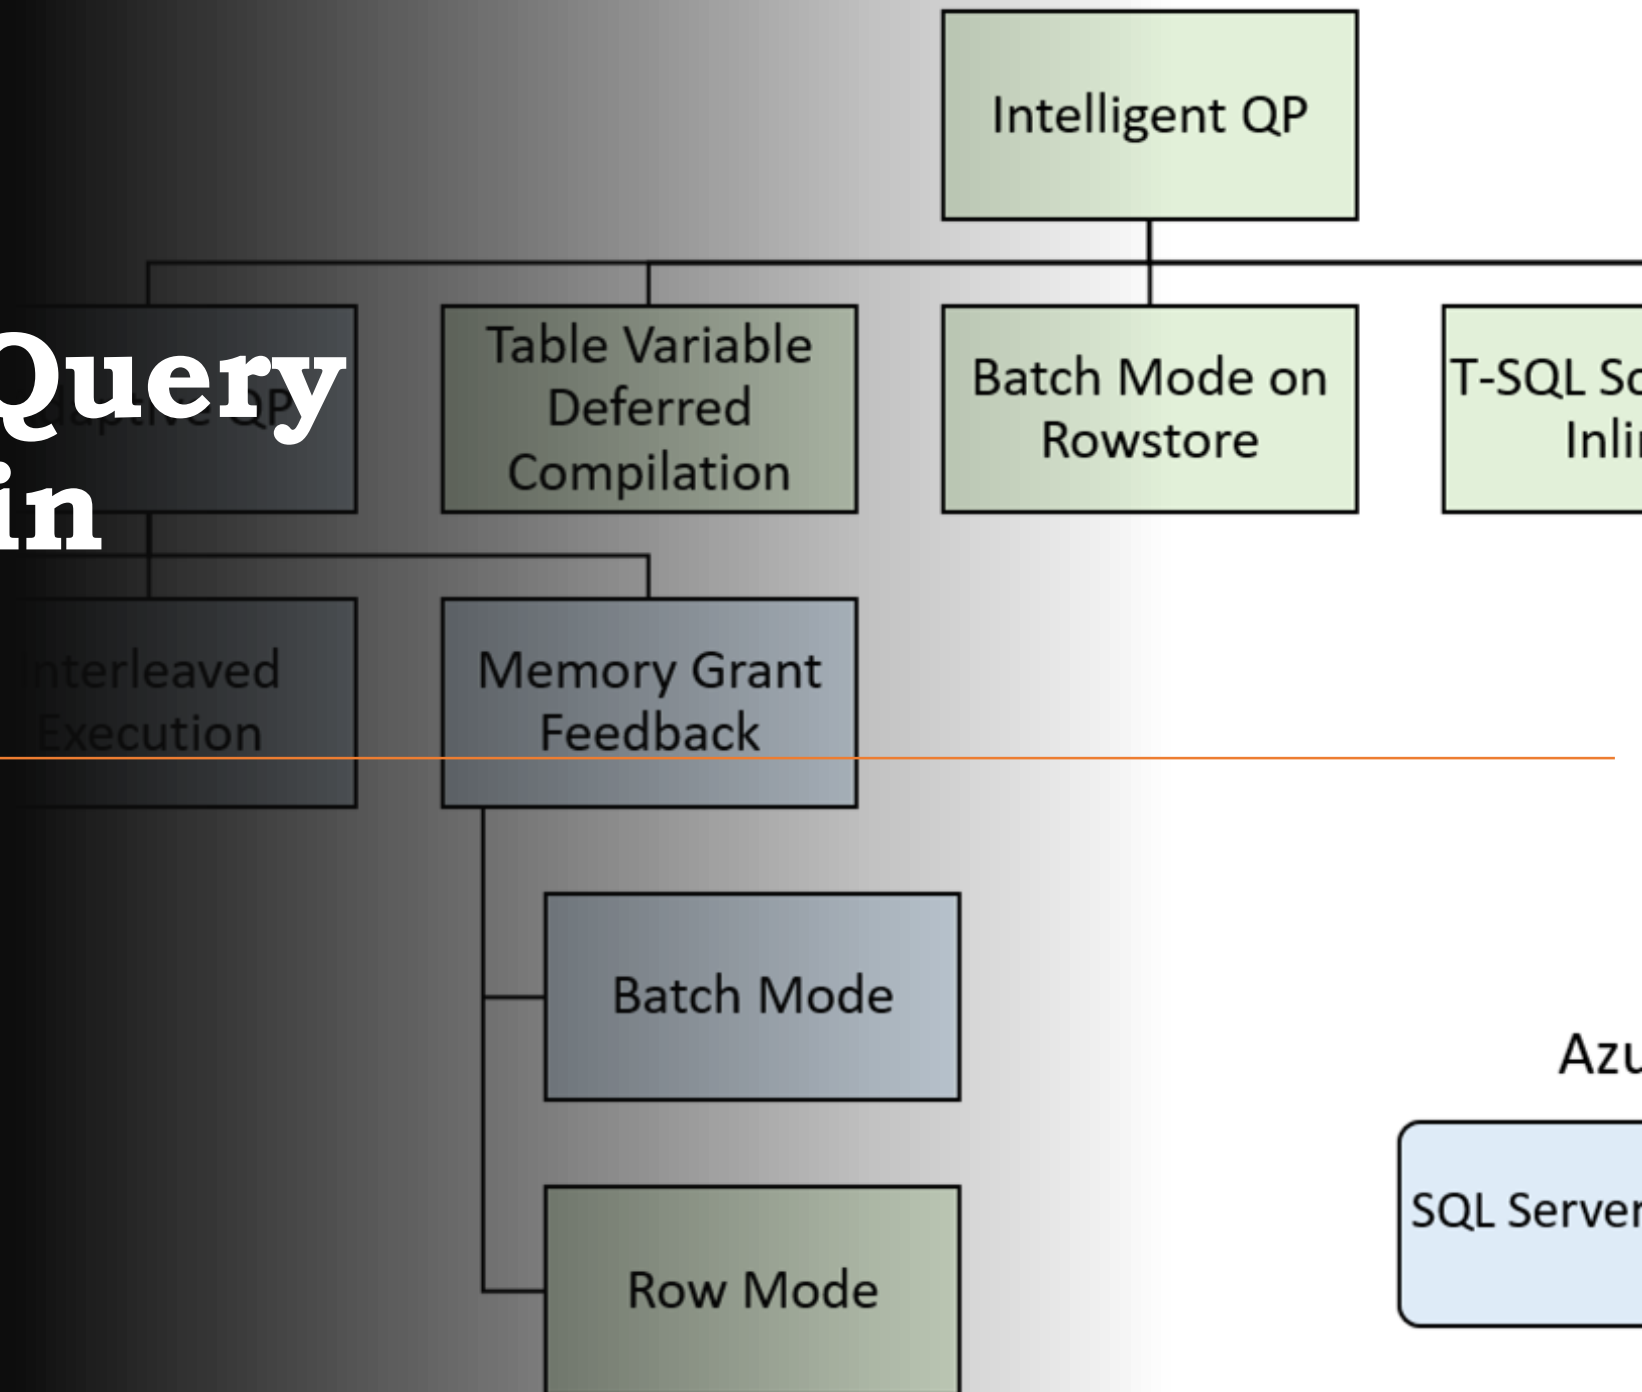


Intelligent Query Processing in SQL Server

Taib Ali





Data Solutions Manager, GMO LLC



<http://sqlworldwide.com/>



/sqlworldwide



@sqlworldwide



Taiob@sqlworldwide.com

Data Professional

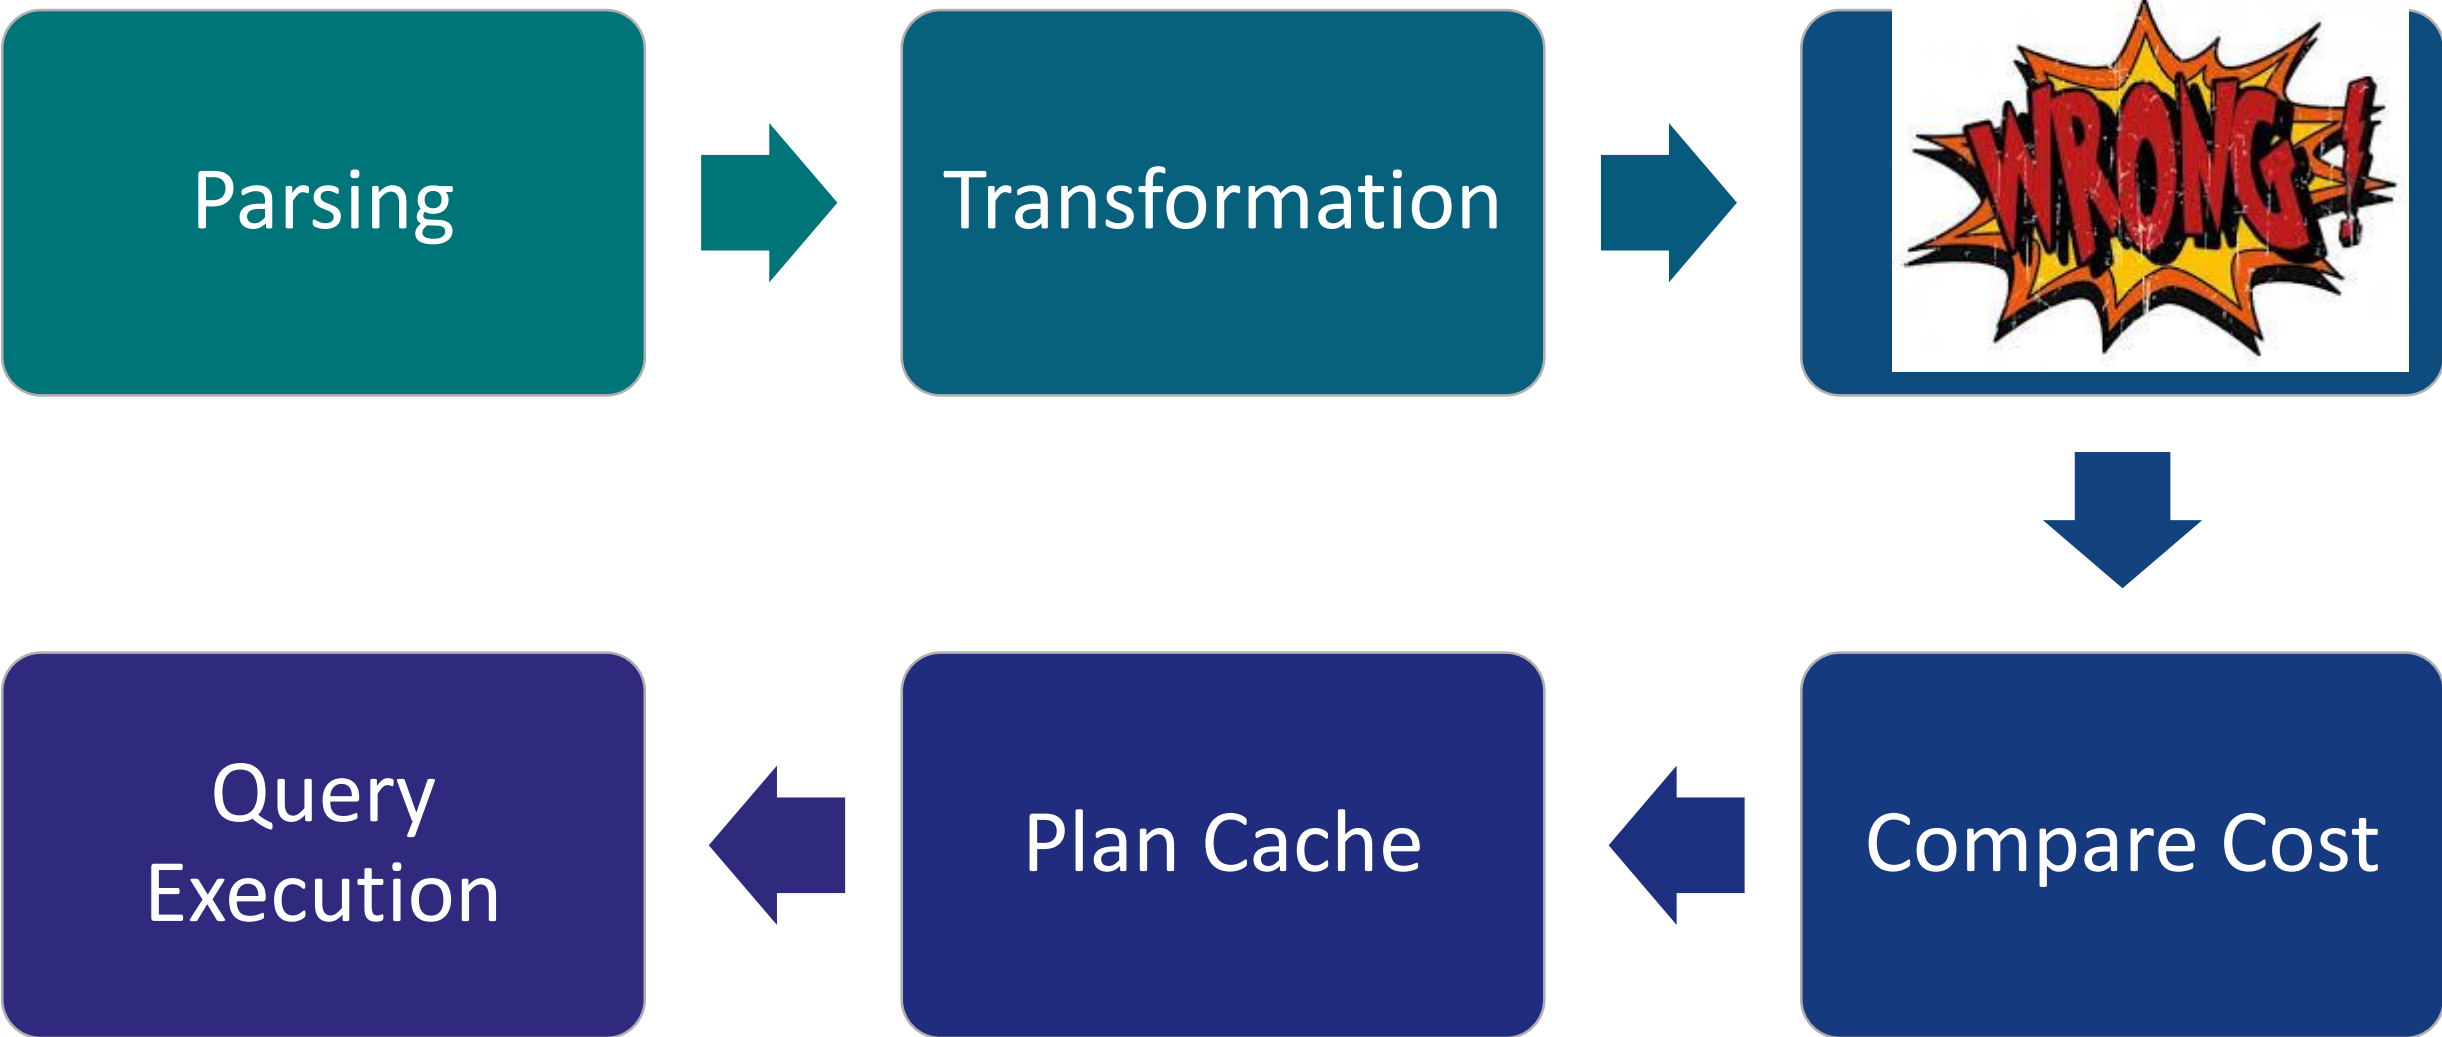
Microsoft Data Platform MVP. 14 Years working with Microsoft Data Platform. Microsoft and MongoDB certified. Worked in ecommerce, healthcare and finance industry.

Giving Back

Board member NESQL user group and PASS DBA virtual group. Organizer of Boston SQL Saturday. Frequent speaker at local and virtual user groups, SQL Saturdays, PASS Summit and Azure events.

When Not Working

Running – One 26.2 and many 13.1 miles. Learning US history. Shuttling 3 kids.



Cost

Parallel

Serial

Memory Grant

In
Memory

Spill to
Disk

Access Method

Seek

Scan

Seek +
Scan

Algorithm

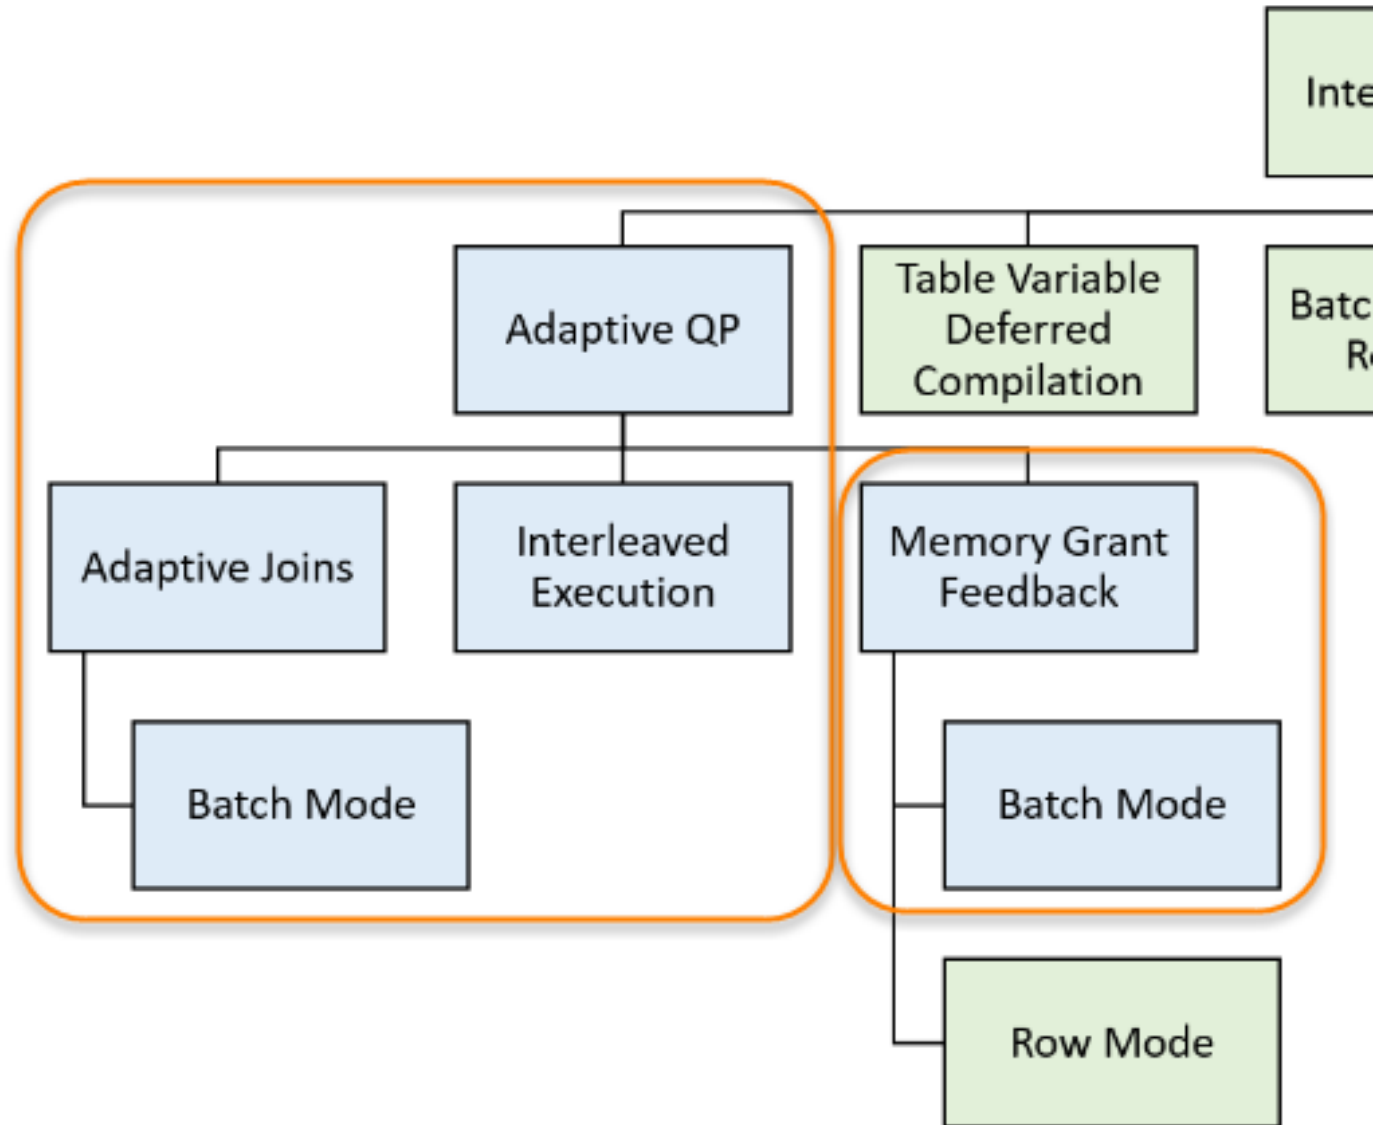
Join

Aggregate

Sort

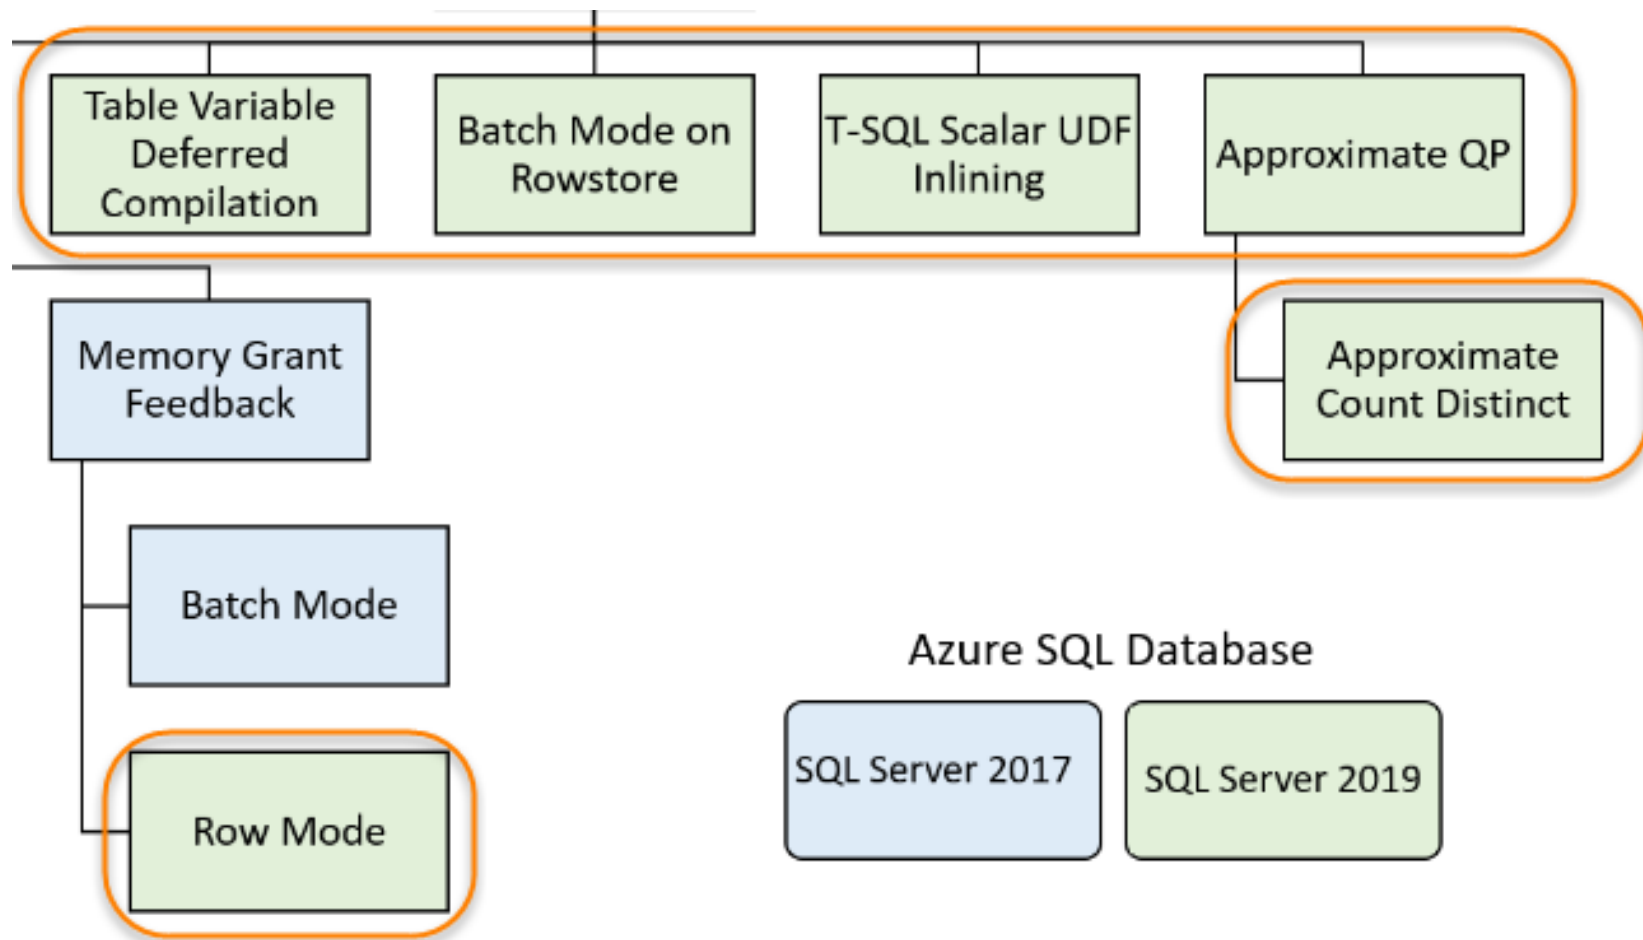


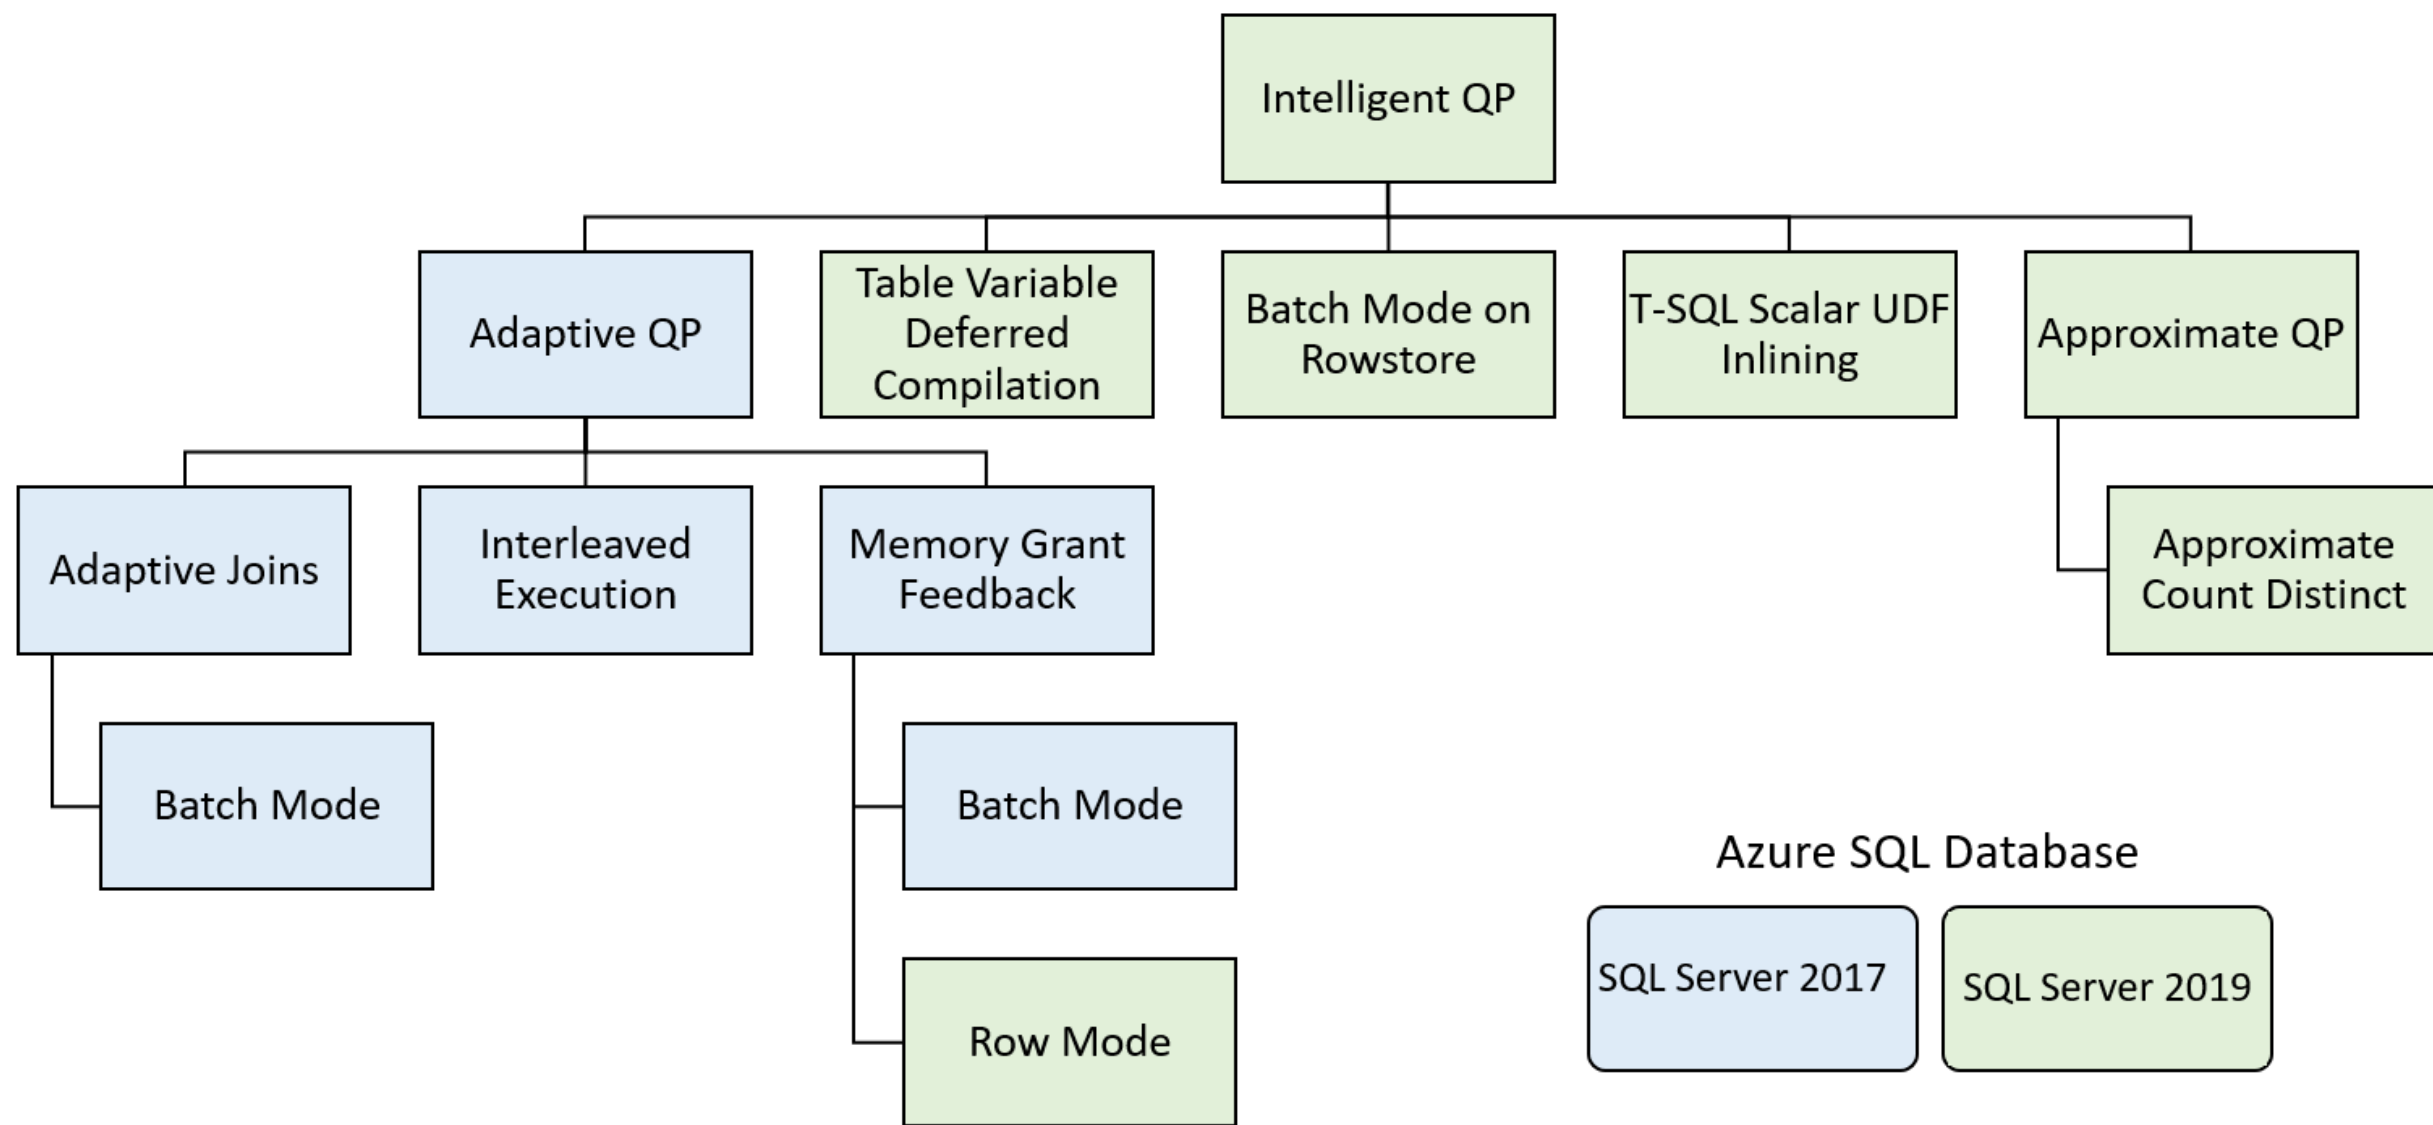
2017





2019







2021

Query Store Hints

- 2021, Preview, Azure SQL Database
- Add hints when you do not have access to the code
- Using query_id (from Query Store)
- Example
 - Recompile on each execution
 - Limit max degree of parallelism for specific queries
 - Disable **optimizer row goal**



2021

CE Feedback

- 2021, Preview, Azure SQL Database
- Significant model estimation errors
- Only repeating workload will qualify
 - Cache persistent plan
- During conflict hard-coded hint will prevail
- Only verified feedback will prevail
- Example
 - Correlation Analysis
 - Row Goal
 - Containment type

SYS.DATABASE_SCOPED_ CONFIGURATIONS

**SYS.DM_EXEC_VALID_US
E_HINTS**



Problem



Solution



Caution



Adaptive Joins Batch Mode

140
150


- Join Hint
- Parameter sensitive query



Adaptive Joins Batch Mode

140
150

- Dynamically switch to a better join
- Nested loop or Hash Join
- Based on threshold



https://docs.microsoft.com/en-us/sql/relational-databases/performance/media/6_aqpjointhreshold.png?view=sql-server-ver15

Adaptive Joins Batch Mode

140
150

- The query is a SELECT statement
- Join needs to be eligible with both Hash and Nested Loop join
- Hash join uses batch mode
- Both join should have same outer reference





Table Variable Deferred Compilation

150

- Works ok with low number of rows but not as rows increase
- Table variables do not have statistics
- Table variables do not have 'Automatic stats creation'
- Only inline index definitions
- Does not trigger recompile
- Fixed cardinality guess of 1



Table Variable Deferred Compilation

150

- Optimizer delays the compilation
 - Same as what temporary table does today
- Accurate cardinality – better execution plan
 - Example Hash join instead of Nested loop join

Table Variable Deferred Compilation

150

- Does not change any other characteristics
- Does not increase recompilation frequency
- Does not fix Parament Sniffing issues





Interleaved Execution MSTVFs

140
150

- MSTVFs have a fixed cardinality guess of
 - 100 in SQL Server 2014 (12.x)
 - 1 in earlier versions



Interleaved Execution MSTVFs

140
150

- Actual row counts are used to make better-informed decision
- Greater performance impact with higher skew



Interleaved Execution MSTVFs

140
150

- Must be read-only and NOT part of a data modification
- Must use [runtime constant](#)



Memory Grant Feedback Batch Mode

140
150

- Performance suffers from incorrect Memory Grant
- Insufficient grant
 - Spill to disk
- Excessive grants
 - Wasted memory
 - Reduced concurrency



Memory Grant Feedback Batch Mode

140
150

- Trigger recalculate
 - Result in a spill to disk
 - Granted memory > 2 x size of the actual used memory
- New SSMS property 'IsMemoryGrantFeedbackAdjusted' to track feedback

Memory Grant Feedback Batch Mode

140
150

- Will disable itself for parameter sensitive queries
- Grants under 1 MB will not be recalculated
- Changes are currently not captured in the Query Store
- Memory Granted honors limitation by the resource governor or query hint





Memory Grant Feedback Row Mode

150

- Row mode memory grant feedback expands on the batch mode



TSQL Scalar UDF Inlining

150

- Iterative invocation
- Lack of costing
- Serial Execution
- Interpreted execution
- Imperative code does not scale



TSQL Scalar UDF Inlining

150

- UDFs are automatically transformed into
 - Scalar Expressions
 - Scalar subqueries
- Further optimization followed by transformation
- Refactors the Imperative code into Relational Algebraic Expression – [Froid Framework](#)
- Resulting execution plan
 - Efficient
 - Set-Oriented
 - Parallel
- New SSMS property
'ContainsInlineScalarTsqlUdfs' to track inlinng



TSQL Scalar UDF Inlining

150

- [Requirements to be eligible](#)
- [Scalar UDF Inlining issues in SQL Server 2019](#)
- [sys.sql_modules](#) – Can this UDF be inlined?
- Can disable within function definition

```
-- Transact-SQL Function Clauses
<function_option>::=
{
    [ ENCRYPTION ]
  | [ SCHEMABINDING ]
  | [ RETURNS NULL ON NULL INPUT | CALLED ON NULL INPUT ]
  | [ EXECUTE_AS_Clause ]
  | [ INLINE = { ON | OFF } ]
}
```

Batch Mode on Rowstore

150

- Row by Row processing is slow and cpu intensive
- Columnstore indexes may not be appropriate for some applications
- Features might restrict use of Columnstore index



Batch Mode on Rowstore

150

- Uses heuristics – during estimation phase
 - Table sizes
 - Operators used
 - Estimated cardinalities
- Additional checkpoints, to evaluate plans with batch mode
- Support for all existing batch mode-enabled operator
- Workload consists of analytics queries especially with joins or aggregates
- Workload that is CPU bound



Batch Mode on Rowstore

150

- Batch mode restriction always applicable
 - Example-Queries involving cursors
- Not applicable for in-memory OLTP tables
- Not applicable for any index other than on-disk heaps and B-trees
- Won't kick in for
 - Large Object (LOB) column
 - XML column
 - Sparse column sets
- Two features are independent





Approximate
Count Distinct

150

- Responsiveness is important than absolute precision
- Example
 - Dashboard scenarios
 - Data science trying to understand data distributions



Approximate Count Distinct

150

- Access of data sets that are millions of rows or higher
- Aggregation of a column or columns that have many distinct values
- Use less memory compare to exhaustive COUNT DISTINCT
- Based on [HyperLogLog](#) algorithm



Approximate
Count Distinct

150

- The function implementation guarantees up to a 2% error rate within a 97% probability

Disabling any of these features without changing the compatibility level

-- SQL Server 2017

```
ALTER DATABASE SCOPED CONFIGURATION SET  
DISABLE_BATCH_MODE_MEMORY_GRANT_FEEDBACK = ON;
```

-- Starting with SQL Server 2019, and in Azure SQL Database

```
ALTER DATABASE SCOPED CONFIGURATION SET  
BATCH_MODE_MEMORY_GRANT_FEEDBACK = OFF;
```

You can also disable any of these features for a specific query by using 'USE HINT' query hint

```
OPTION (USE HINT ('DISABLE_BATCH_MODE_MEMORY_GRANT_FEEDBACK'));
```

Resource

- **Intelligent Query Processing in SQL databases**
- **Intelligent Query Processing Demos**
- **Compatibility Certification**
- **Query Processing Architecture Guide**
- **Get Your Scalar UDFs to Run Faster Without Code Changes**
- **Batch Mode Bitmaps in SQL Server by Paul White**
- **Introducing Batch Mode Adaptive Memory Grant Feedback by Joe Sack**



SQL2019 CU8
SSMS 18.7.1



@sqlworldwide



linkedin.com/in/sqlworldwide



sqlworldwide.com



taio@sqlworldwide.com

