

Think Like The Cardinality Estimator

Taiob Ali

He/Him





Taib Ali

Data Solutions Manager, GMO LLC



<http://sqlworldwide.com/>



/sqlworldwide



@sqlworldwide



taib@sqlworldwide.com

Data Professional

Microsoft Data Platform MVP. 14 Years working with Microsoft Data Platform. Microsoft and MongoDB certified. Worked in ecommerce, healthcare and finance industry.

Giving Back

Board member NESQL user group and PASS DBA virtual group. Organizer of Boston SQL Saturday. Frequent speaker at local and virtual user groups, SQL Saturdays, PASS Summit and Azure events.

When Not Working

Running – 1x26.2 and 30+x 13.1, Learning US history. Shuttling 3 kids.

Agenda

- Definition
 - What is Cardinality?
 - Why Cardinality Matters?
 - DBCC SHOW_STATISTICS
 - Magic Numbers
-

Definition

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
70	2035	330	123	3	110
71	2043	333	90	3	111
72	2051	317	94	3	105.6667
73	2055	106	114	1	106
74	2061	249	122	2	124.5
75	2065	107	108	1	107
76	2073	326	111	3	108.6667
77	2077	127	129	1	127
78	2083	234	119	2	117
79	2091	332	106	3	110.6667
80	2095	118	135	1	118
81	2103	338	125	3	112.6667
82	2107	112	110	1	112
83	2113	224	109	2	112
84	2121	375	118	3	125
85	2127	224	101	2	112
82	5151	554	101	5	115
84	5151	352	118	3	152
83	5113	554	108	5	115
85	5101	115	110	1	115
81	5103	338	152	3	115.6667

Predicate

- **Expression** = TRUE, FALSE, UNKNOWN
 - Join
 - Filter
 - Where
 - Having

```
SELECT
    cus.CustomerID,
    COUNT(0) AS [NumOfOrders]
FROM
    sales.Orders AS ord
JOIN
    sales.Customers AS cus
ON
    ord.CustomerID=cus.CustomerID
WHERE
    ord.OrderDate='2013-01-01'
GROUP BY cus.CustomerID
HAVING COUNT(0) > 2 ;
GO
```

Join Predicate

Filter Predicate
Where/Having

Predicate Selectivity

Fraction of rows from the input set of the predicate that satisfy the predicate

$$\frac{[\text{\# rows that pass the predicate}]}{[\text{total number of rows}]}$$

Predicate Selectivity

```
SELECT  
    COUNT(0) AS [NumOfOrders]  
FROM  
    sales.Orders  
WHERE  
    CustomerID=577;  
GO
```

75 rows for customerID (577)

73595 rows for all customer

$75 / 73595 = .0010$

Density

How often duplicate values occur in a column

$$\frac{1}{[\text{\# of distinct values in a column}]}$$

```
SELECT
    COUNT(DISTINCT customerID) AS [DistinctCusId]
FROM
    sales.Orders;
GO
```

663 distinct customerID

$$1 / 663 = 0.0015083$$

What is Cardinality Estimation?

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROW
70	2035	330	123	3
71	2043	333	90	3
72	2051	317	94	3
73	2055	106	114	1
74	2061	249	122	2
75	2065	107	108	1
76	2073	326	111	3
77	2077	127	129	1
78	2083	234	119	2
79	2091	332	106	3
80	2095	118	135	1
81	2103	338	125	3
82	2107	112	110	1
83	2113	224	109	2
84	2121	375	118	3
85	2127	224	101	2
82	5153	554	101	5
84	5151	352	118	3
83	5113	554	108	5
85	5103	115		

Structured Query Language

Cardinality estimation (CE) is the process by which the Query Optimizer derives the estimated number of rows for a query plan

Clustered Index Seek (Clustered)

Scanning a particular range of rows from a clustered index.

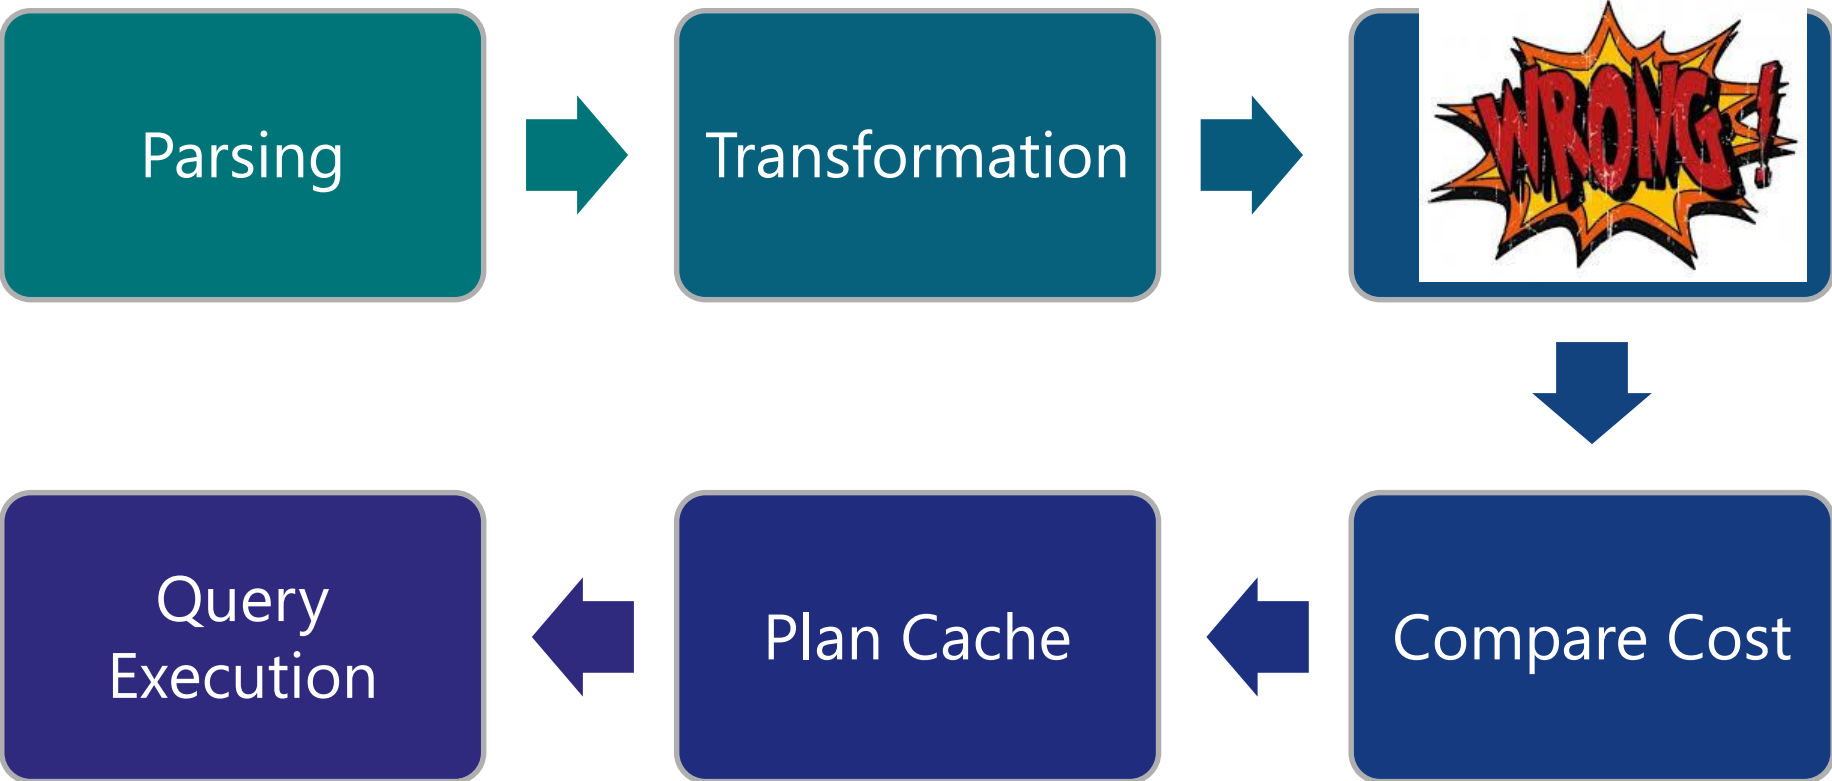
Physical Operation	Clustered Index Seek
Logical Operation	Clustered Index Seek
Actual Execution Mode	Row
Estimated Execution Mode	Row
Storage	RowStore
Actual Number of Rows	12
Actual Number of Batches	0
Estimated Operator Cost	0.0032904 (8%)
Estimated I/O Cost	0.003125
Estimated Subtree Cost	0.0032904
Estimated CPU Cost	0.0001654
Estimated Number of Executions	1
Number of Executions	1
Estimated Number of Rows	7.61577
Estimated Row Size	143 B

Actual

Estimated

Why Cardinality Matters?

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
70	2035	330	123	3	110
71	2043	333	90	3	111
72	2051	317	94	3	105.6667
73	2055	106	114	1	106
74	2061	249	122	2	124.5
75	2065	107	108	1	107
76	2073	326	111	3	108.6667
77	2077	127	129	1	127
78	2083	234	119	2	117
79	2091	332	106	3	110.6667
80	2095	118	135	1	118
81	2103	338	125	3	112.6667
82	2107	112	110	1	112
83	2113	224	109	2	112
84	2121	375	118	3	125
85	2127	224	101	2	112
82	S1S3	554	101	5	115
84	S1S1	352	118	3	152
83	S1I3	554	108	5	115
85	S1O3	115	110	1	115



Cost

Parallel

Serial

Memory Grant

In
Memory

Spill to
Disk

Access Method

Seek

Scan


Seek +
Scan

Algorithm

Join

Aggregate

Sort



<http://www.historyofmurphyslaw.com>

Missing Statistics

Stale Statistics

Inadequate Statistics Sample Rate

Maximum 200 Steps in Histogram

Parameter Sniffing

Out-of-Model Query Constructs

DBCC SHOW_STATISTICS

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS
70	2035	330	123	3
71	2043	333	90	3
72	2051	317	94	3
73	2055	106	114	1
74	2061	249	122	2
75	2065	107	108	1
76	2073	326	111	3
77	2077	127	129	1
78	2083	234	119	2
79	2091	332	106	3
80	2095	118	135	1
81	2103	338	125	3
82	2107	112	110	1
83	2113	224	109	2
84	2121	375	118	3
85	2127	224	101	2

Header

- **Meta data about the statistics.**

Density Vector

- **How many unique values are present within a column or columns?**

Histogram

- **Frequency of data within the first key column of the statistics.**

STAT_HEADER

Name	Updated	Rows	Rows Sampled	Steps
NCI_FilteredContactPersonID	Mar 31 2017 4:55PM	50299	50299	200

Deprecated

Density	Average key length	String Index	Filter Expression	Unfiltered Rows
0.00945746	8	NO	[[contactpersonid]>(2000)]	73595

DENSITY_VECTOR

1/Number of
distinct values in
column

Names of columns
in the prefix

All density	Average Length	Columns
0.002164502	4	ContactPersonID
1.988111E-05	8	ContactPersonID, OrderID

Histogram

	RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
70	2035	330	123	3	110
71	2043	333	90	3	111
72	2051	317	94	3	105.6667
73	2055	106	114	1	106
74	2061	249	122	2	124.5
75	2065	107	108	1	107
76	2073	326	111	3	108.6667
77	2077	127	129	1	127
78	2083	234	119	2	117
79	2091	332	106	3	110.6667
80	2095	118	135	1	118
81	2103	338	125	3	112.6667
82	2107	112	110	1	112
83	2113	224	109	2	112
84	2121	375	118	3	125
85	2127	224	101	2	112

RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
2083	334	119	2	117
2091	332	106	3	110.6667

RANGE_HI_KEY	RANGE_ROWS	EQ_ROWS	DISTINCT_RANGE_ROWS	AVG_RANGE_ROWS
2083		119		
Between 2084 and 2090	332		3	110.66
2091		106		



- Single Predicate
 - Histogram direct hit
 - Histogram intra step
 - Scaling
 - Distinct



- Multiple Predicates
 - Conjunction
 - Disjunction
- Parameter Sniffing
- Unknown
- Ascending Key

DEMO

SQL2019 CU8
SSMS 18.7.1



Reference

- [Statistics](#)
- [Query Tuning Fundamentals](#)
- [DBCC SHOW STATISTICS \(Transact-SQL\)](#)
- [13 Things You Should Know About Statistics and the Query Optimizer](#)
- [Cardinality Estimation for Multiple Predicates](#)
- [New Trace Flag to Fix Table Variable Performance](#)
- [Ascending key Issue TF 2389 and 2390](#)
- [Optimizing Query Plans with the SQL Server 2014 Cardinality Estimator](#)
- [Cardinality Estimation \(SQL Server\)](#)



@sqlworldwide



linkedin.com/in/sqlworldwide



sqlworldwide.com



taio@sqlworldwide.com

